


# Practical Data Science and Machine Learning with GPUs

Accelerate. Connect. Empower.

Stanley Seibert  
Director of Community Innovation

# Part 1: Machine Learning Basics in Anaconda


# What is Machine Learning?



# What can machine learning help us do?


- **Predict the future:**
  - ✓ *"How much corn will I sell next month?"*
- **Reveal hidden information:**
  - ✓ *"How many truck tires likely fail to meet our durability specification based on the temperature data collected during their manufacture?"*
- **Identify structure in large data sets:**
  - ✓ *"Can business loan applications be grouped by common attributes?"*
- **Find unusual trends:**
  - ✓ *"Which items in my grocery store have seen an abnormal sales decrease outside of historical seasonal variation?"*

# ML Categories: Regression



# ML Categories: Classification


airplane



automobile



bird



cat



deer




dog



frog



horse



ship




truck



- Assign each sample to one or more categories.
- Categories known ahead of time
- Can be **exclusive** or **inclusive**


# ML Categories: Clustering

- Divide a set of items into groups
- Categories are **not** known ahead of time



[http://scikit-learn.org/stable/auto\\_examples/cluster/plot\\_dbscan.html](http://scikit-learn.org/stable/auto_examples/cluster/plot_dbscan.html)

# ML Categories: Anomaly Detection






# ML Categories: Generation



<http://genekogan.com/works/style-transfer/>

# Deep Learning

## Inception v3 Network



# of parameters: 23.8M  
# of layers: ~140


# Critical ML Questions, Part 1

- What is the purpose of the model I am trying to create?
- What specific, quantitative model output will help me achieve that purpose?
- Are there publications, blog posts, conference talks describing a similar use case?
- What relevant data is available for this model? Can more be gathered?
- Is there additional domain knowledge that would be helpful for selecting training data?
- Are there experts in the data I can talk to?

# Critical ML Questions, Part 2

- Are there security or ethical concerns associated with the training data and the future uses of this model?
- What accuracy has already been achieved in our organization, and how much improvement is required for success?
- Can the business benefits of improved accuracy be quantified?
- How will I know the model is working once deployed?


# Today's (Toy) Problem



- Can we classify small color images into one of 10 categories?
- Goal is better than 75% accuracy
- Will use Keras (w/ TensorFlow) and GPU acceleration for training



# Hardware: NVIDIA GPUs



Tesla K80  
4992 CUDA cores (total)  
4.29 TFLOPS (single precision)  
480 GB/sec



× 26

# Software: Anaconda Distribution



## ANACONDA NAVIGATOR

Desktop Portal to Data Science

## ANACONDA PROJECT

Portable Data Science Encapsulation

## DATA SCIENCE LIBRARIES

Data Science IDEs




Analytics & Scientific Computing



Visualization



Machine Learning



and many more!


## CONDA

Data Science Packages & Environment Manager

GPU Accelerated Packages:

- TensorFlow (+ Keras)
- Theano
- PyTorch
- Caffe
- Numba

# Data: CIFAR10 Image Set

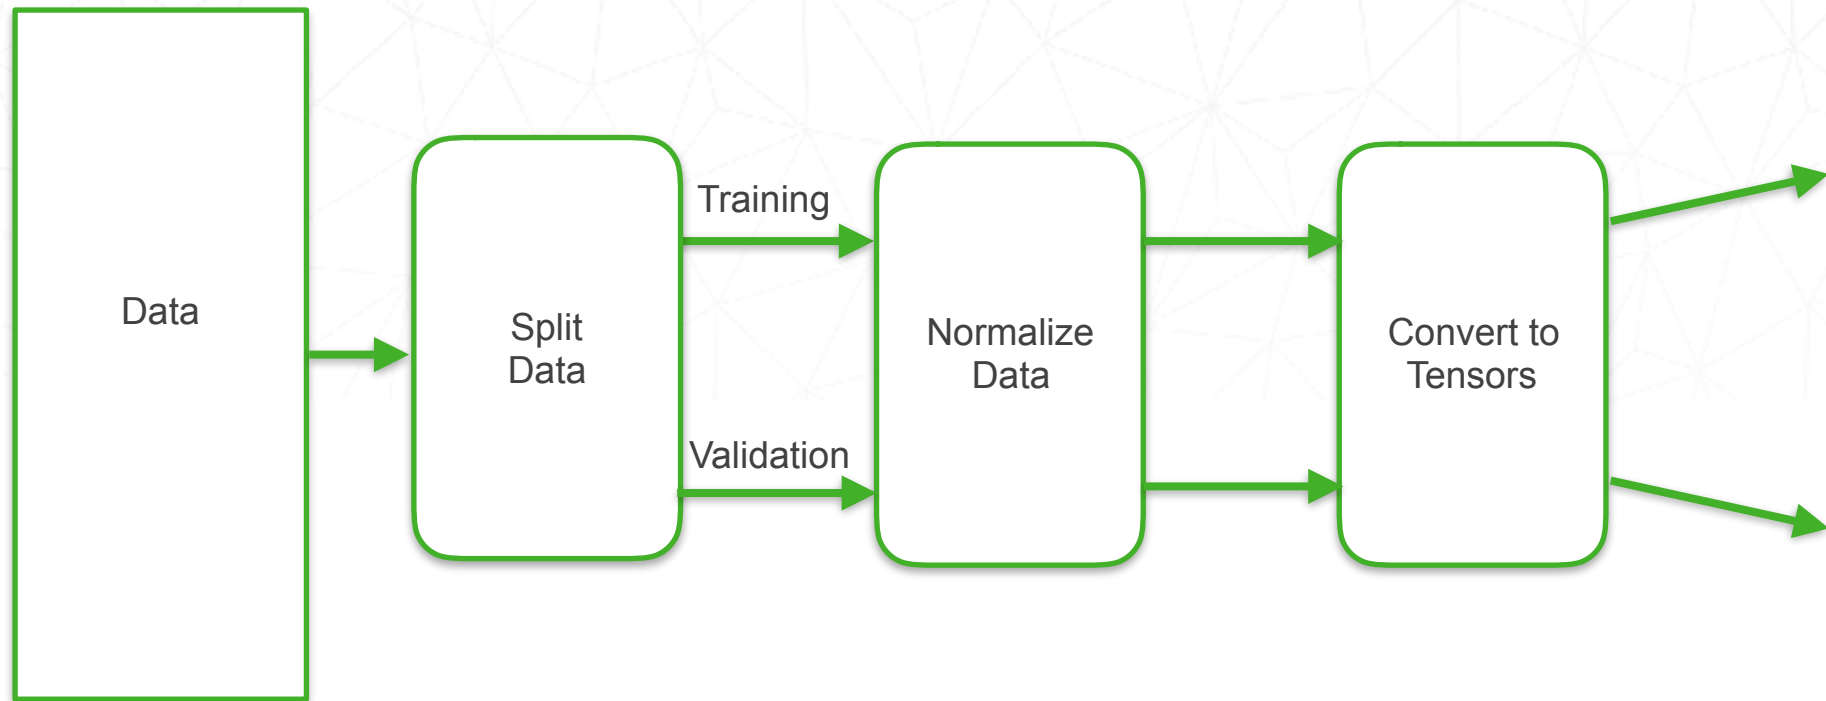


- Classic data set for benchmarking image classification
  - 60000 images
  - 32x32 color pixels
  - 10 classes
  - 6000 images from each class
  - Classes are mutually exclusive

<https://www.cs.toronto.edu/~kriz/cifar.html>



# Data Preparation



# One-Hot Encoding of Categories

*Category #*

1
0
2
3
3
1
2



0	1	2	3
0	1	0	0
1	0	0	0
0	0	1	0
0	0	0	1
0	0	0	1
0	1	0	0
0	0	1	0

# Fooling Yourself With Data



Non-representative Training Data



Non-stationary distribution


# Exercise 1: Data Exploration and Prep

# Part 2: Introduction to Deep Learning with Keras


# What is a neural network?

- *A biologically-inspired method for making models out of a large number of very simple computational operations connected to form a large mesh*
- Very flexible.
- Require a lot of data to train.
- Although described in physical terms (nodes, layers, etc), they are always implemented in software using array math.

# Node







# Layer






# Activation Functions



# Network



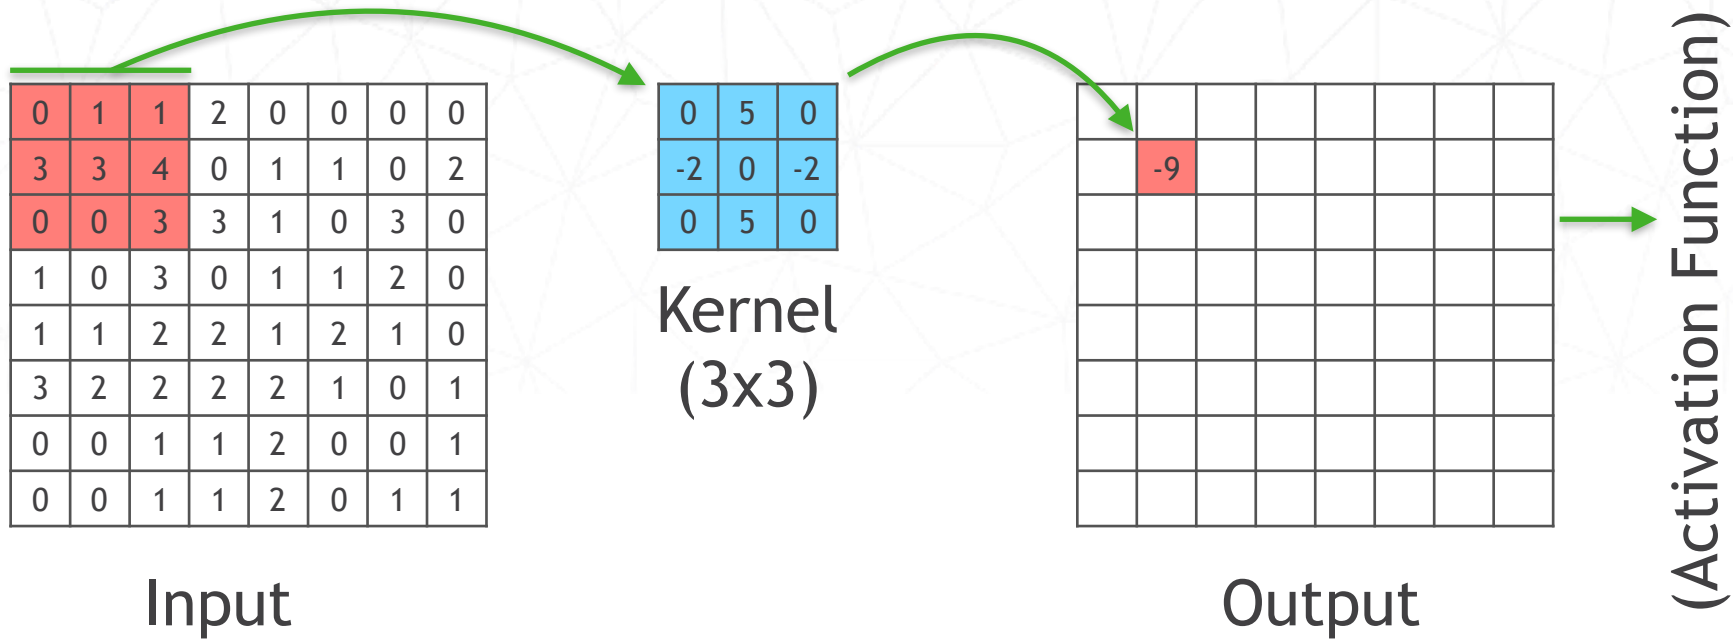
# Layer Types

- "Dense" or "Fully connected"
- Convolutional
- Pooling
- LSTM
- (output) Softmax
- (output) Sigmoid

# Dense Layers




# Convolutional Layers



# Pooling Layers

4	2	1	0
3	3	4	0
0	0	3	3
1	0	3	0



4	4
1	3

# Sigmoid vs. Softmax Output

If you set a threshold, can accommodate non-exclusive classes and "none of the above" results

airplane	0.010
automobile	0.110
bird	0.005
cat	0.160
deer	0.003
dog	0.007
frog	0.150
horse	0.010
ship	0.100
truck	0.005

*Sigmoid*

Membership in a class is determined by having the largest value


(Output always sums to 1)

airplane	0.017
automobile	0.196
bird	0.009
cat	0.286
deer	0.005
dog	0.013
frog	0.268
horse	0.018
ship	0.179
truck	0.009

*Softmax*

# An Image Classification Network


VGG19



<https://www.slideshare.net/ckmarkohchang/applied-deep-learning-1103-convolutional-neural-networks>




# Training a Network



# Loss Functions

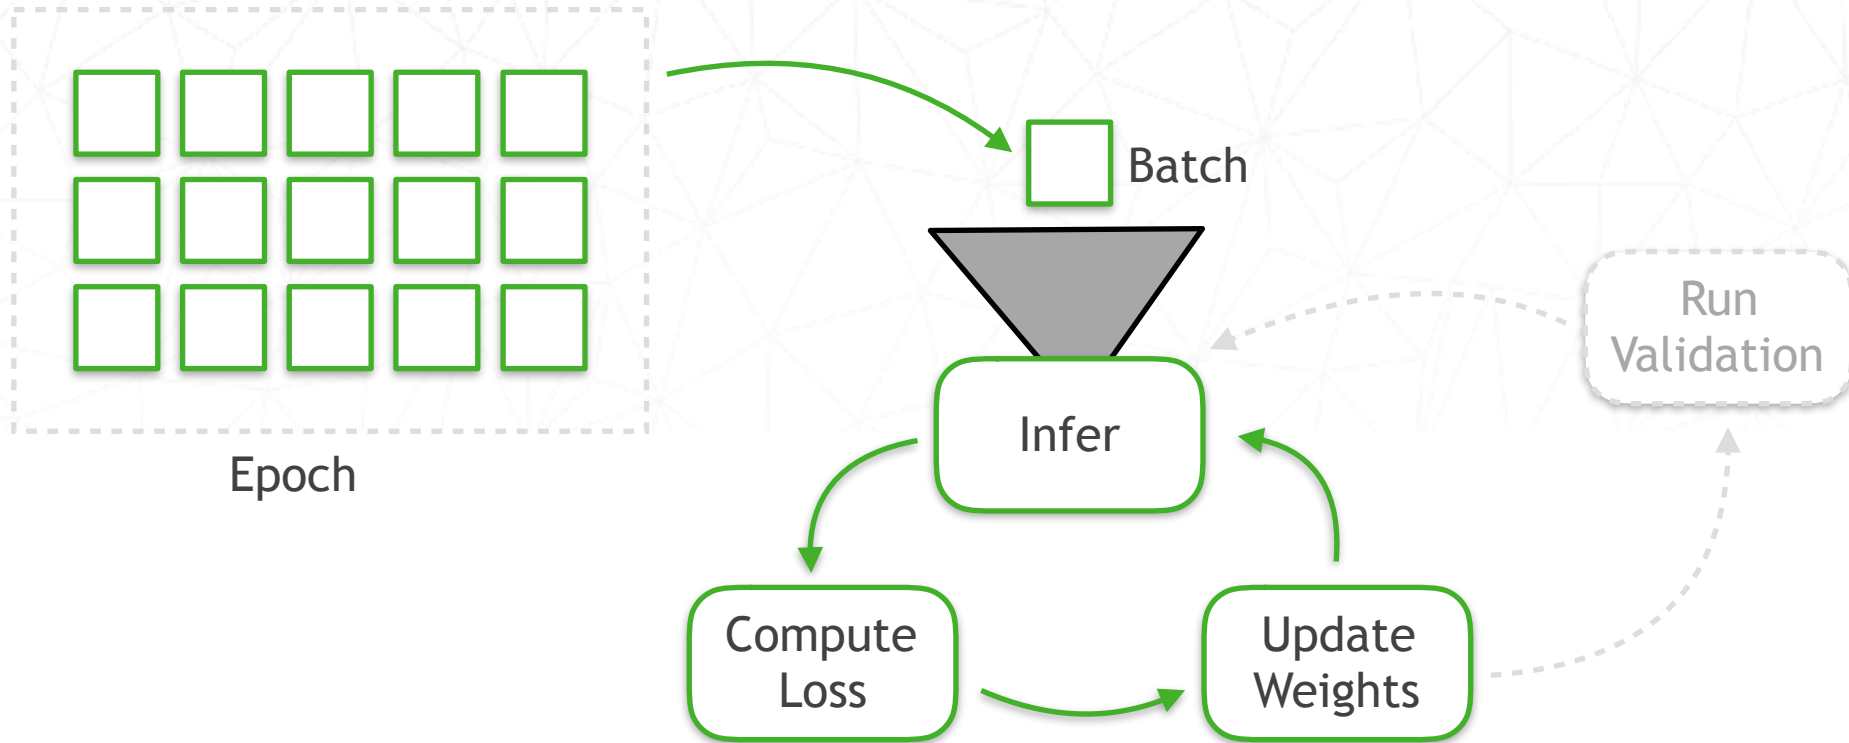
- The optimizer will try to minimize the loss function on the training data
- *Picking a loss function depends on the kind of model you are building*
- Some good default choices:
  - Regression: `mean_squared_error`
  - Binary classification: `binary_crossentropy`
  - Multi-category classification: `categorical_crossentropy`

# Backpropagation & Optimization



- Work backwards through network to compute each weight affects loss function
- Optimization algorithm adjusts each weight according to some strategy.


# Batches and Epochs



# Exercise 2: Deep Learning with Keras


# Part 3: Evaluating Models and Troubleshooting

# Most important question in ML:



*How do I know when I am done?*

# Overfitting




Gap caused  
by model  
overfitting



# What to do about overfitting?

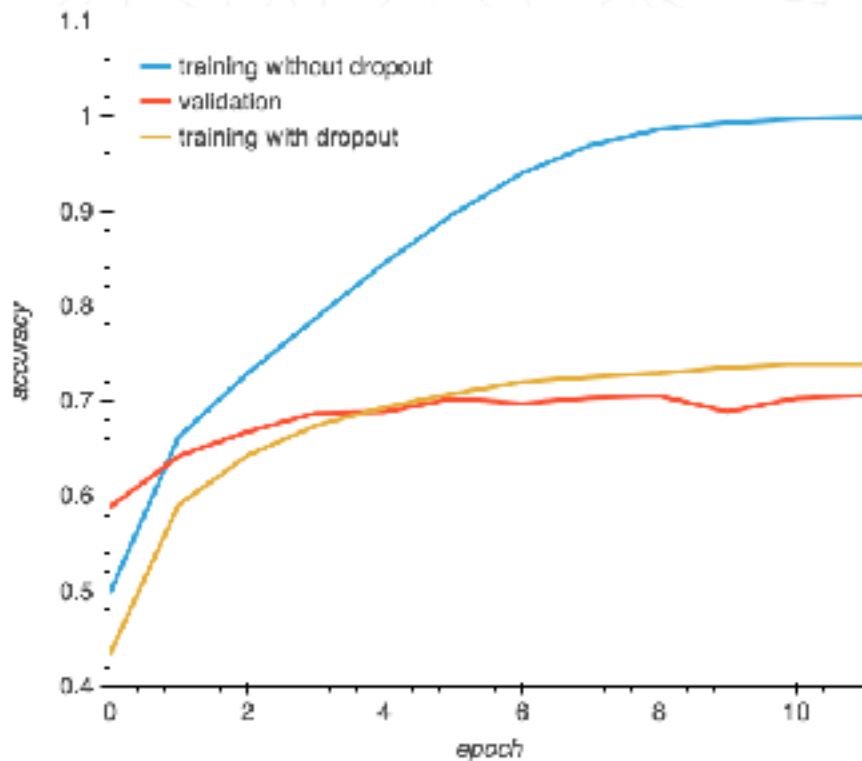
- **Overfitting does not mean your model is bad.**
  - Most models will overfit after enough training epochs
  - *The validation data accuracy is what you care about*
- Techniques exist to control overfitting:
  - Regularization
  - Dropout layers
  - Reduce size of network
  - Get more data

# Dropout Layers



- Only active during training
- Fraction is an adjustable parameter
  - *Usually 0.25 to 0.5*


# After Adding Dropout



Much less  
overfitting

But now we have a  
new problem...

# Underfitting

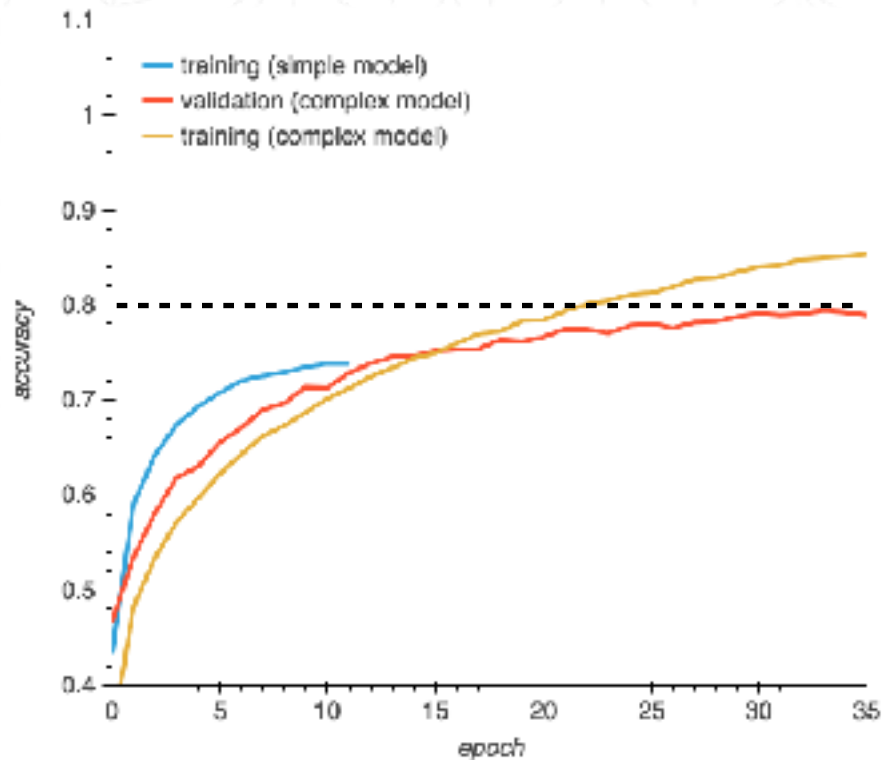


Model accuracy never reaches the goal.

Possible causes:

- Model too simple?
- Not enough training data?
- Mislabeled data?
- Optimizer learning rate?
- ...

# Fixing Underfitting



Changed 2 things

- Increased model complexity (added extra convolutional layers)
- Changed optimizer algorithm


Trial and error is still the norm...

# Limitations of Accuracy



- Accuracy is like the "check engine light"
- Can hide some problems
- *Interpretation is hard if your training data is unbalanced*
- Important to look at other measures:
  - False positive rate
  - False negative rates
  - Confusion matrix
  - Inspect fail cases!

# Confusion Matrix

True values



# Effect of Batch Size



*8 training epochs*



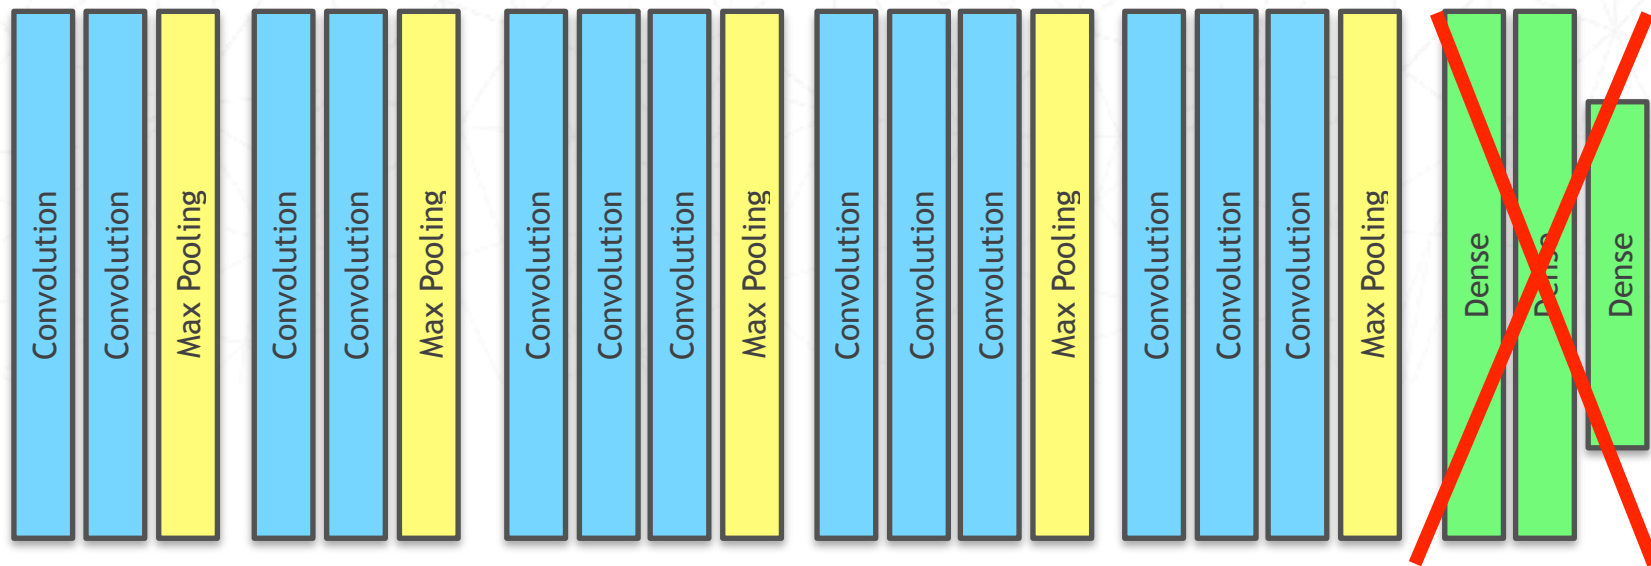
# Hyperparameters

- Choices that affect the structure or training of the network, but are not optimized during training.
- Examples include:
  - Number of nodes in each layer
  - Batch size
  - Dropout rates
- Encapsulate your training code into a Python function that you can call many times for different hyperparameter choices.
- Be wary of overfitting! Hold aside a test dataset to check at the very end.

# Exercise 3: Evaluating Models

# Part 4: Advanced Topics and Deployment


# Building on Existing Models



VGG16 Pretrained on ImageNet

Chop top layers off and  
retrain your own layers

# Curriculum Learning




- Train the model in multiple passes
- Use increasingly hard examples on each pass
- Easy if your data has a "difficulty" knob

*Bengio, et al, ICML 2009*

# Using Multiple GPUs


```
keras.utils.training_utils.multi_gpu_model()
```



# Getting to Deployment

- This was the point of doing all this, right?
- Tools and best practices evolving rapidly
- Questions to consider:
  - What hardware is available in production?
  - Bulk processing or online processing? (or both?)
  - What is the performance requirement (throughput, latency)?
  - Does the model need optimization? (memory size or speed)
  - How should I package it, along with dependencies?

# Determining Hardware Needs



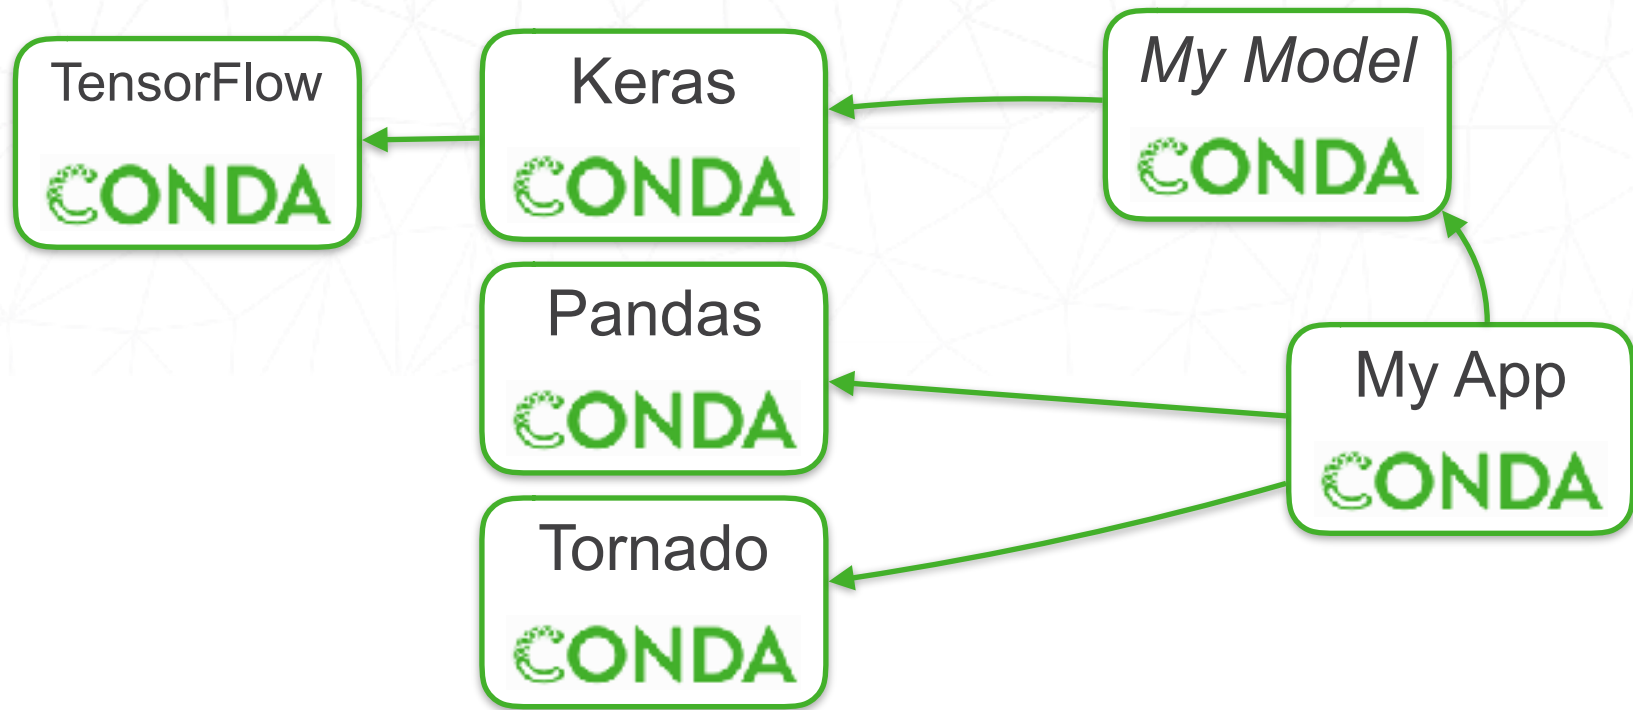
Inference in batches,  
or one at a time?

Do you need a GPU in  
production?

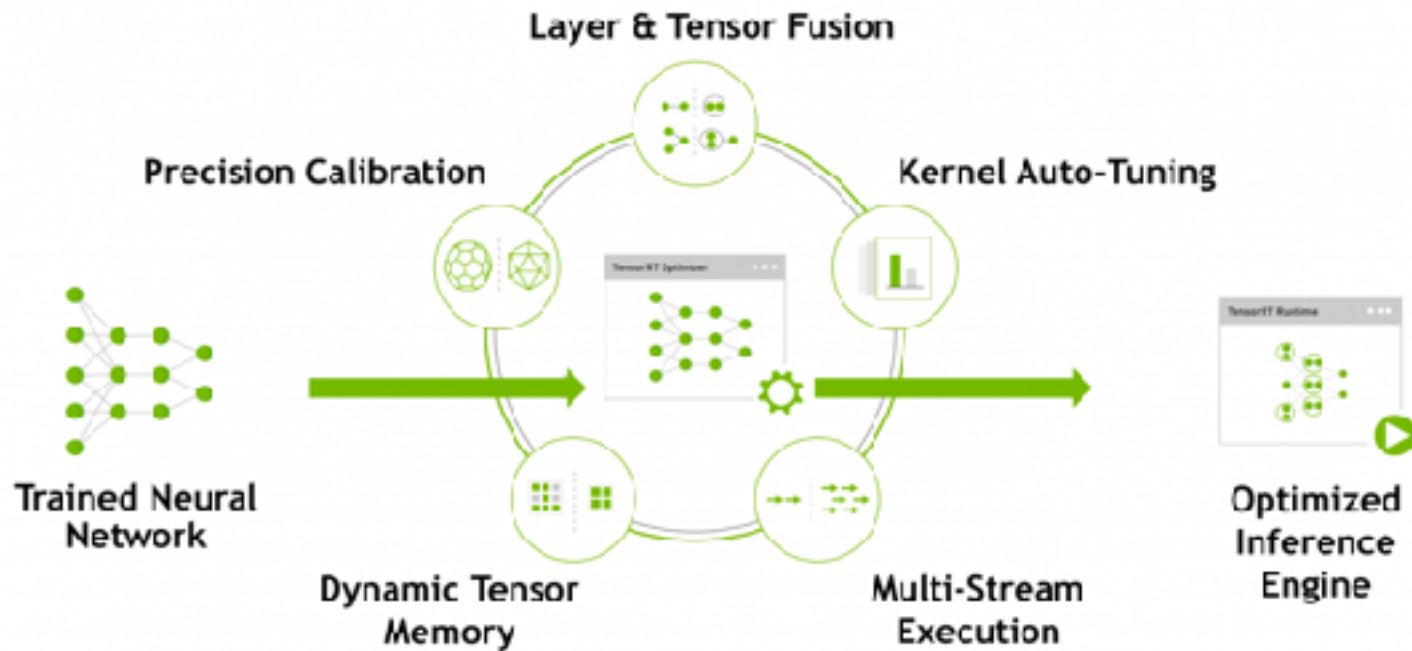
Smaller GPUs may  
also be sufficient



# Packing Models with Conda



# TensorRT



<https://developer.nvidia.com/tensorrt>

# A Data Scientist's Job Is Never Done

- Important to monitor deployed models:
  - Does accuracy on new data match training/validation data?
  - Snapshot incorrect predictions for study
  - Keep expanding your training data
  - *Make sure any data collection complies with your security and privacy policies*
- Version your models, just like software!

# Exercise 4: Loading/Saving/ Benchmarking Models

# Conclusion

# Summary

- Getting started with GPU-accelerated ML has never been easier!
- Remember the process:
  1. Define your problem
  2. Identify your dataset
  3. Design a network
  4. Train the network
  5. Check your work and iterate as needed
  6. Package for production
  7. Monitor deployed models for effectiveness

# Further Reading

- Tutorial Notebooks:  
<https://github.com/ContinuumIO/ac2018-dl-gpu>
- *Deep Learning with Python*, by François Chollet  
<https://www.manning.com/books/deep-learning-with-python>
- Neural Network Playground (experiment in your browser!)  
<http://playground.tensorflow.org/>
- Keras Documentation  
<https://keras.io/>
- Want code examples? Google search for "Keras [use case]".  
Examples:  
["keras image segmentation"](#)